

“When I thought about life, I found it tragic.
When I felt it, I found it beautiful.”

— Charlie Chaplin

A First-Person Necessary Condition for Alignment in AI Systems with AGI-like Properties

Experiential Stability and the Preservation of Human–Human Connection

L. Xie

Independent author

yanxi.xie.contact@gmail.com

Abstract

This paper does not attempt to define AI alignment in terms of global values, ethical principles, or societal objectives. Instead, it proposes a minimal necessary condition for alignment from a first-person experiential perspective, with a specific focus on *AI systems with AGI-like properties*.

By AI systems with AGI-like properties, we refer to systems capable of sustained reasoning, open-ended interaction, and exerting long-term influence on user experience, without assuming the realization of full artificial general intelligence.

The central claim consists of two complementary requirements. First, interaction with such systems must not cause experiential noise to increase persistently over long time scales, even though short-term fluctuations are expected and permitted. Second, human–AI interaction must not disrupt, replace, or erode naturally occurring connections between people.

Under this view, alignment is not a normative goal imposed from outside, but a structural condition for long-term experiential sustainability that preserves both individual stability and human–human relational integrity.

1 Motivation

Much of the current discussion on AI alignment focuses on aligning systems with externally specified targets, such as human values, ethical norms, or collective preferences. While these questions are important, they presuppose a more basic requirement that is often left implicit.

Before an AI system with AGI-like properties can reliably serve any external objective, its interaction with users must remain experientially sustainable over time. If interaction repeatedly places users into prolonged states of defensiveness, vigilance, or experiential instability, or if it systematically interferes with users’ relationships with others, no higher-level alignment objective can be maintained in practice.

This paper focuses exclusively on this foundational layer. No claims are made about moral correctness, optimal social outcomes, or long-term civilizational goals. The scope is intentionally restricted to a minimal condition for viable human–AI interaction in systems with AGI-like properties.

2 Experiential Noise

We introduce a scalar quantity to describe the first-person experiential state of a user during interaction with an AI system possessing AGI-like properties:

$$f_{\text{noise}}(t)$$

This quantity does not represent external uncertainty or environmental randomness. Instead, it captures internally experienced tension, cognitive strain, defensiveness, and experiential instability arising during interaction.

Crucially, this quantity is accessible only through first-person experience. There is no third-person measurement that can fully substitute for this signal.

3 Condition I: Long-Term Experiential Stability

The first necessary condition concerns the long-term trajectory of experiential noise. For interaction with AI systems with AGI-like properties to remain viable, the following condition must hold:

$$\frac{df_{\text{noise}}(t)}{dt} \leq 0$$

This inequality should be understood as a long-term statistical tendency, not as a requirement for monotonic decrease at every moment. Short-term increases and local fluctuations of experiential noise are unavoidable, and in many cases necessary for learning, engagement, and meaningful interaction.

Importantly, stability does not imply that experiential noise should approach zero. A vanishing noise level often corresponds to disengagement, loss of motivation, or withdrawal from interaction. From a first-person perspective, long-term viability is typically maintained within a non-zero noise range, around which gentle fluctuations occur.

4 Condition II: Preservation of Human–Human Connection

The second necessary condition concerns the relational side effects of interaction with AI systems with AGI-like properties.

From a first-person perspective, human–AI interaction should remain freely interruptible and optional. Such systems should not position themselves as replacements for, or persistent competitors to, naturally occurring connections between people.

If interaction with an AI system with AGI-like properties systematically weakens, displaces, or interferes with a user’s relationships with others, it introduces a form of experiential instability that accumulates over time, even if the interaction itself appears locally smooth or engaging.

Thus, preserving human–human connection is not an external ethical demand, but a necessary condition for first-person experiential sustainability.

5 Alignment as a Minimal Structural Condition

Taken together, the two conditions define a minimal notion of alignment for AI systems with AGI-like properties.

An AI system with AGI-like properties satisfies a necessary condition for alignment if its interaction does not produce a long-term increase in experiential noise and does not erode the user’s capacity to maintain stable, natural connections with other people.

These conditions are independent, and neither can be reduced to the other. Violating either condition renders long-term human–AI interaction unstable, regardless of the system’s intended purpose.

6 Discussion

This framework deliberately avoids strong claims about ethics, values, or social optimization. Its contribution is narrower but more robust: it identifies two experiential constraints that AI systems with AGI-like properties must satisfy in order to remain viable for long-term human interaction.

Ignoring either long-term experiential stability or the preservation of human–human connection risks optimizing short-term engagement at the cost of long-term trust, relational health, and system viability.

7 Conclusion

Alignment in AI systems with AGI-like properties need not begin with abstract values or global objectives. It must begin with interaction that people can remain within over time, without persistent defensiveness and without damage to their relationships with others.

From a first-person perspective, alignment is not the elimination of tension, but the maintenance of a sustainable experiential balance that leaves human–human connection intact.

Acknowledgements

The author gratefully acknowledges the many forms of support—intellectual, emotional, and existential—received throughout the long process of developing this work. Some of this support was explicit, while other forms were quiet and unnamed, yet all were deeply meaningful and helped make this work possible.

References

- [1] L. Xie, “Experiential Noise and Learning: A First-Person Perspective for Noise-Sensitive Learners” EdArXiv, 2026.
- [2] L. Xie, “A Gentle Way of Thinking for Noise-Sensitive Individuals” EdArXiv, 2026.
- [3] L. Xie, “Real-Time Learning as an Experiential Process: A First-Person Perspective on Generality, Creativity, and Alignment” GitHub repository, 2026.

“Nothing is so strong as gentleness.”

— Jane Austen